



Operating Systems

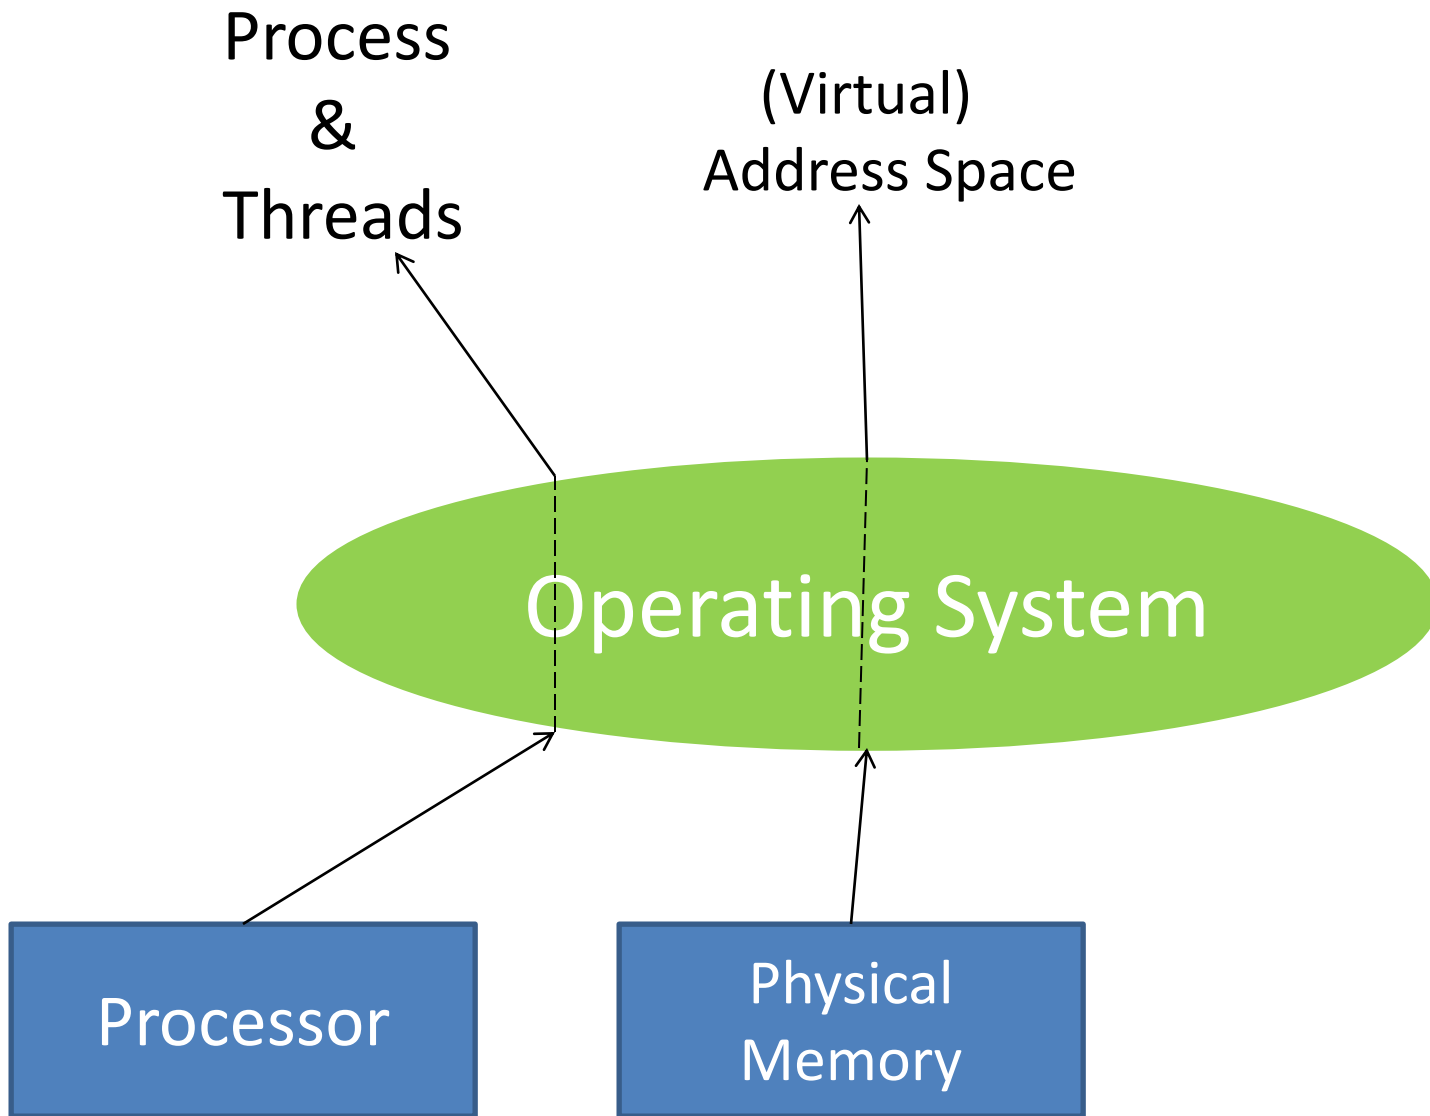
File Systems

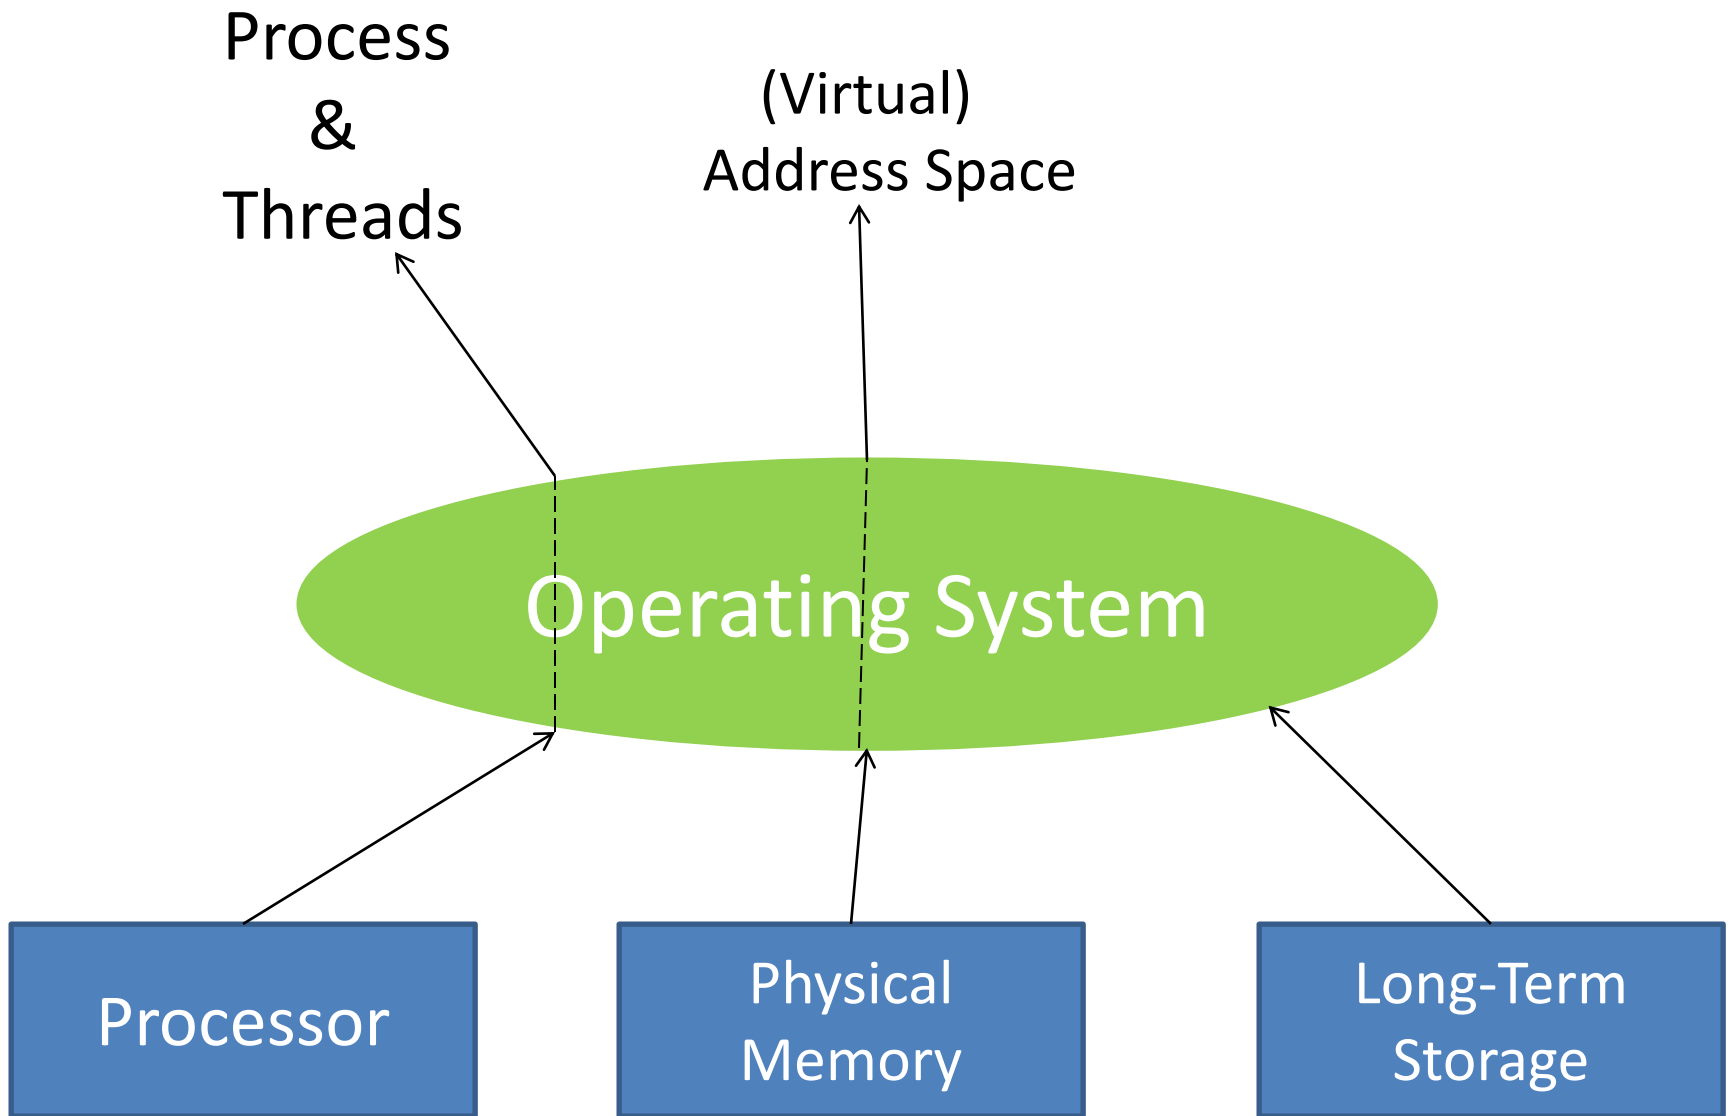
Mohamed Zahran (aka Z)

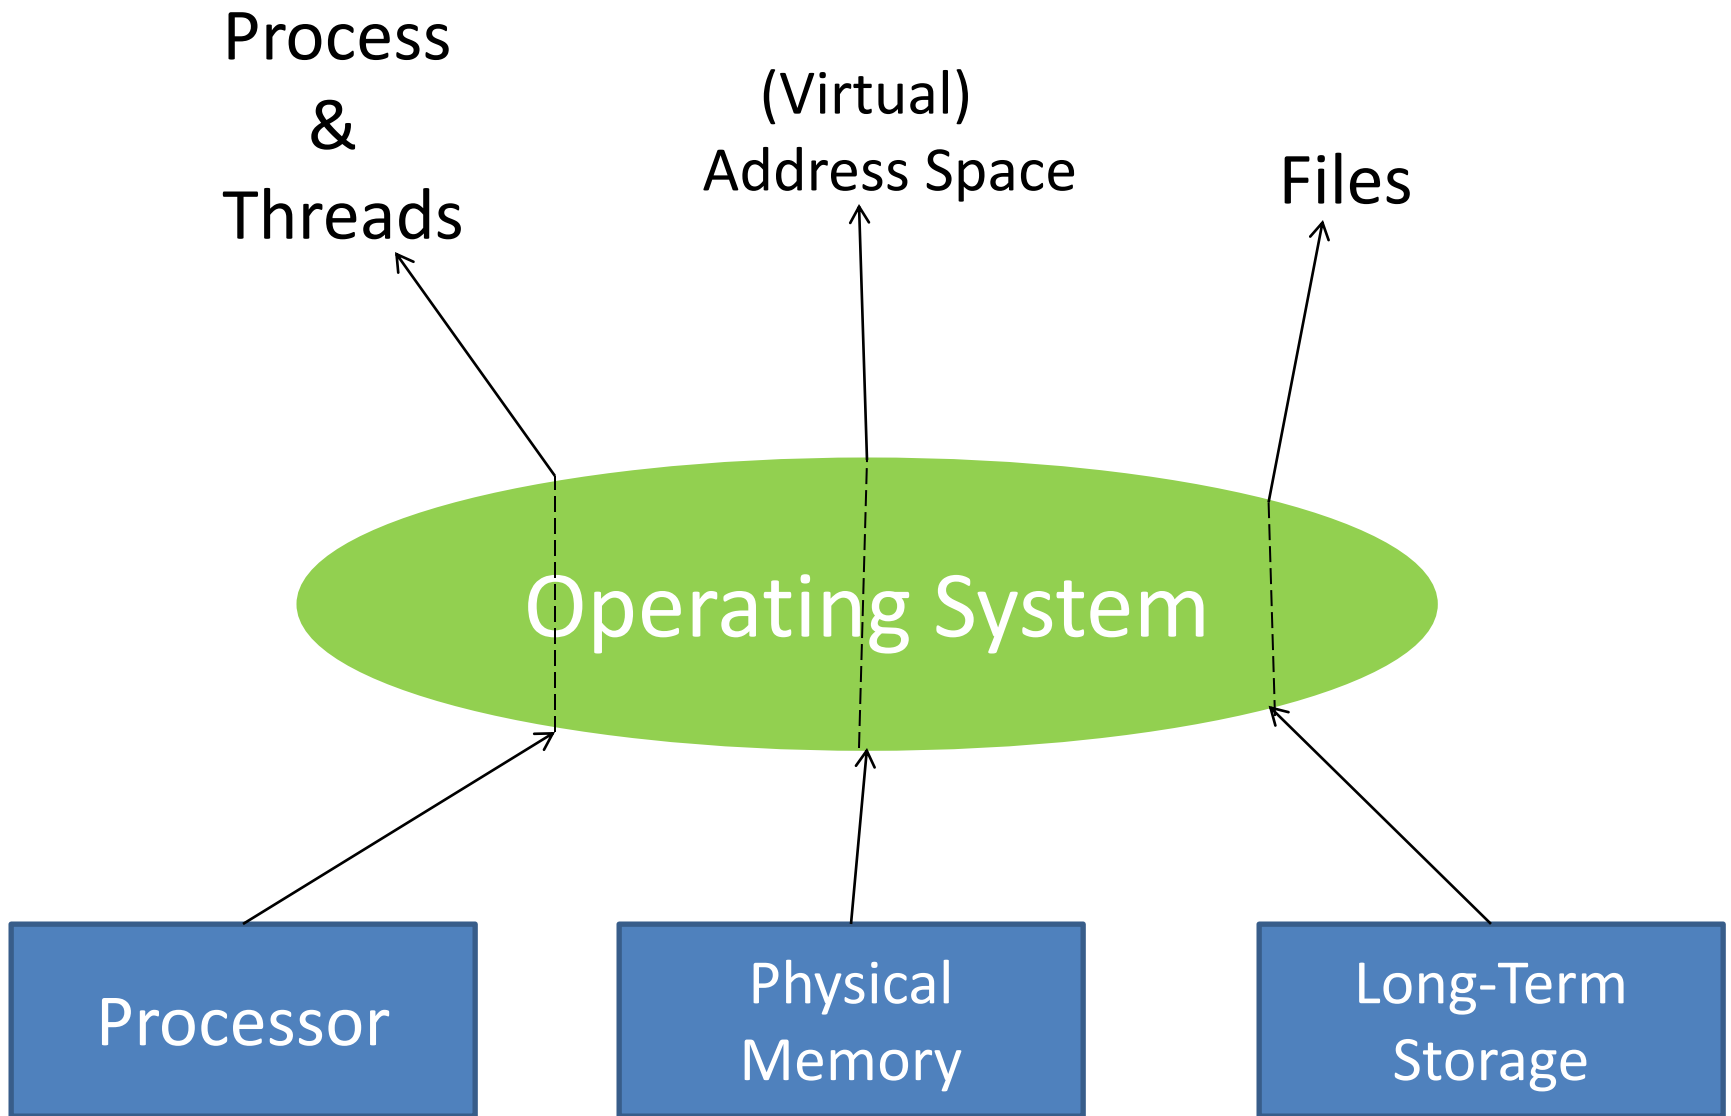
mzahran@cs.nyu.edu

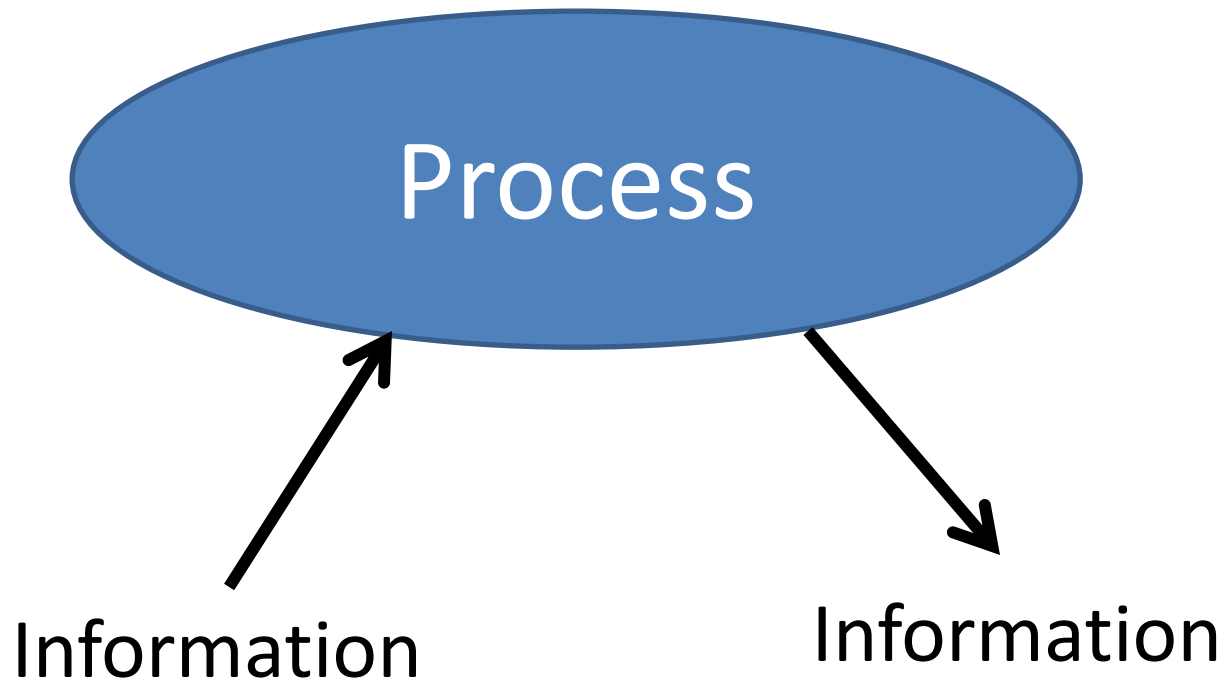
<http://www.mzahran.com>











Is it OK to keep this information only in the process address space?

Shortcomings of Process Address Space


- Virtual address space may not be enough storage for all information.
- Information is lost when the process terminates, is killed, or the computer crashes.
- Multiple processes may need the information at the same time.

Requirements for Long-term Information Storage

- Store very large amount of information.
- **Persistence:** Information must survive the termination of the process using it.
- Multiple processes must be able to access the information concurrently.

Two Key Abstractions for Storage

- **Files**
 - Simply a linear array of bytes
 - In most systems, the OS does not know much about the structure of the file.
- **Directories**



Directories and files have low-level names beside the name used by the user.

Issues with long term storage

- How do you find information?
- How do you keep one user from reading another user's data?
- How do you know which **blocks** are free?



Let's assume a disk consists of a **linear sequence of fixed-size blocks** supporting two operations: read block x and write block x.


FILES




```
graph TD; FILES([FILES]) --> Naming[Naming];
```

Naming


- Shields the user from details of file storage.
- In general, files continue to exist even after the process that creates them terminates.

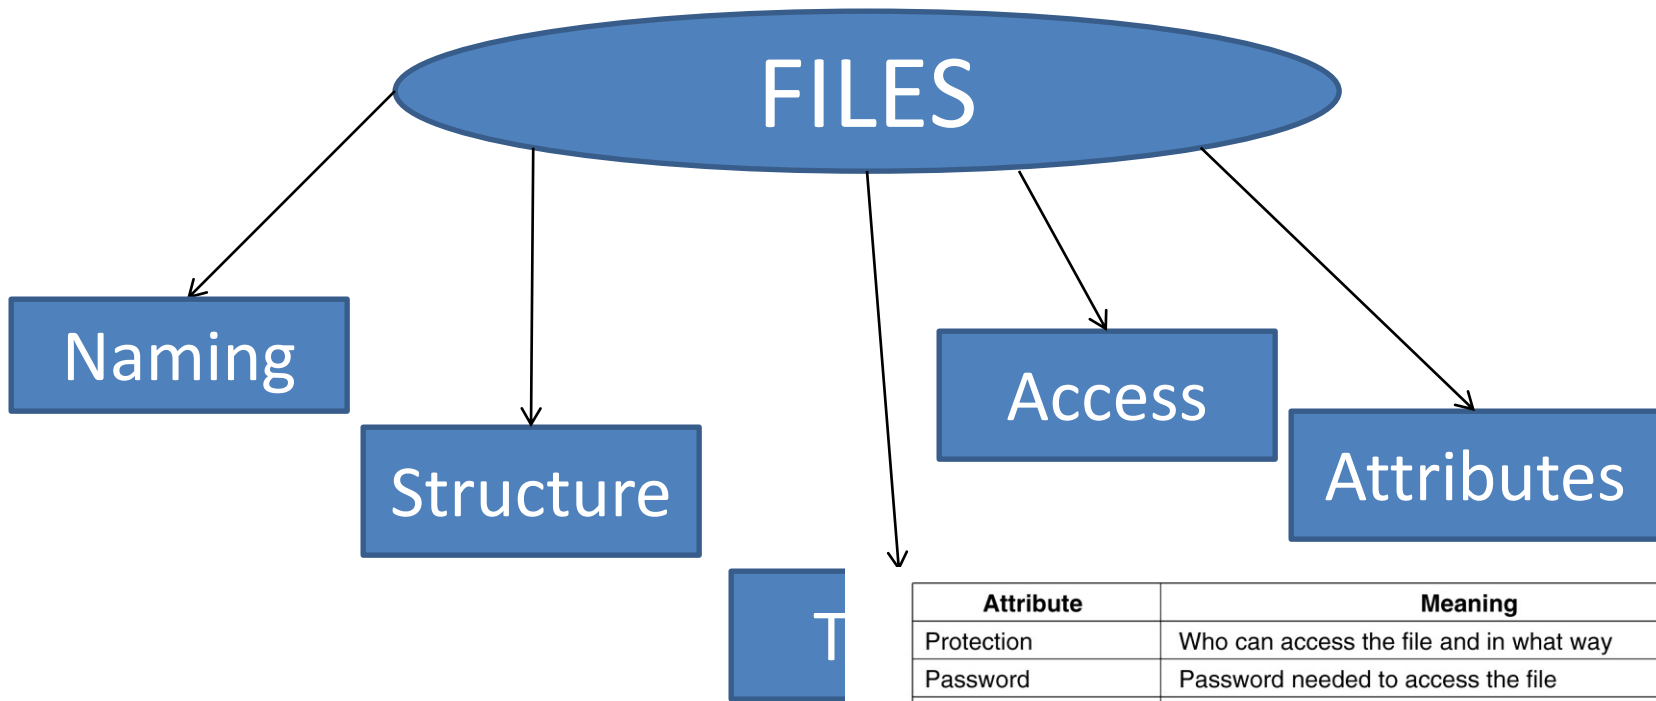


- OS sees it as a linear stream of bytes.
- Users can have different views (e.g. record, ...)
- Programs can impose structures (e.g. jpg, mp3, ...)

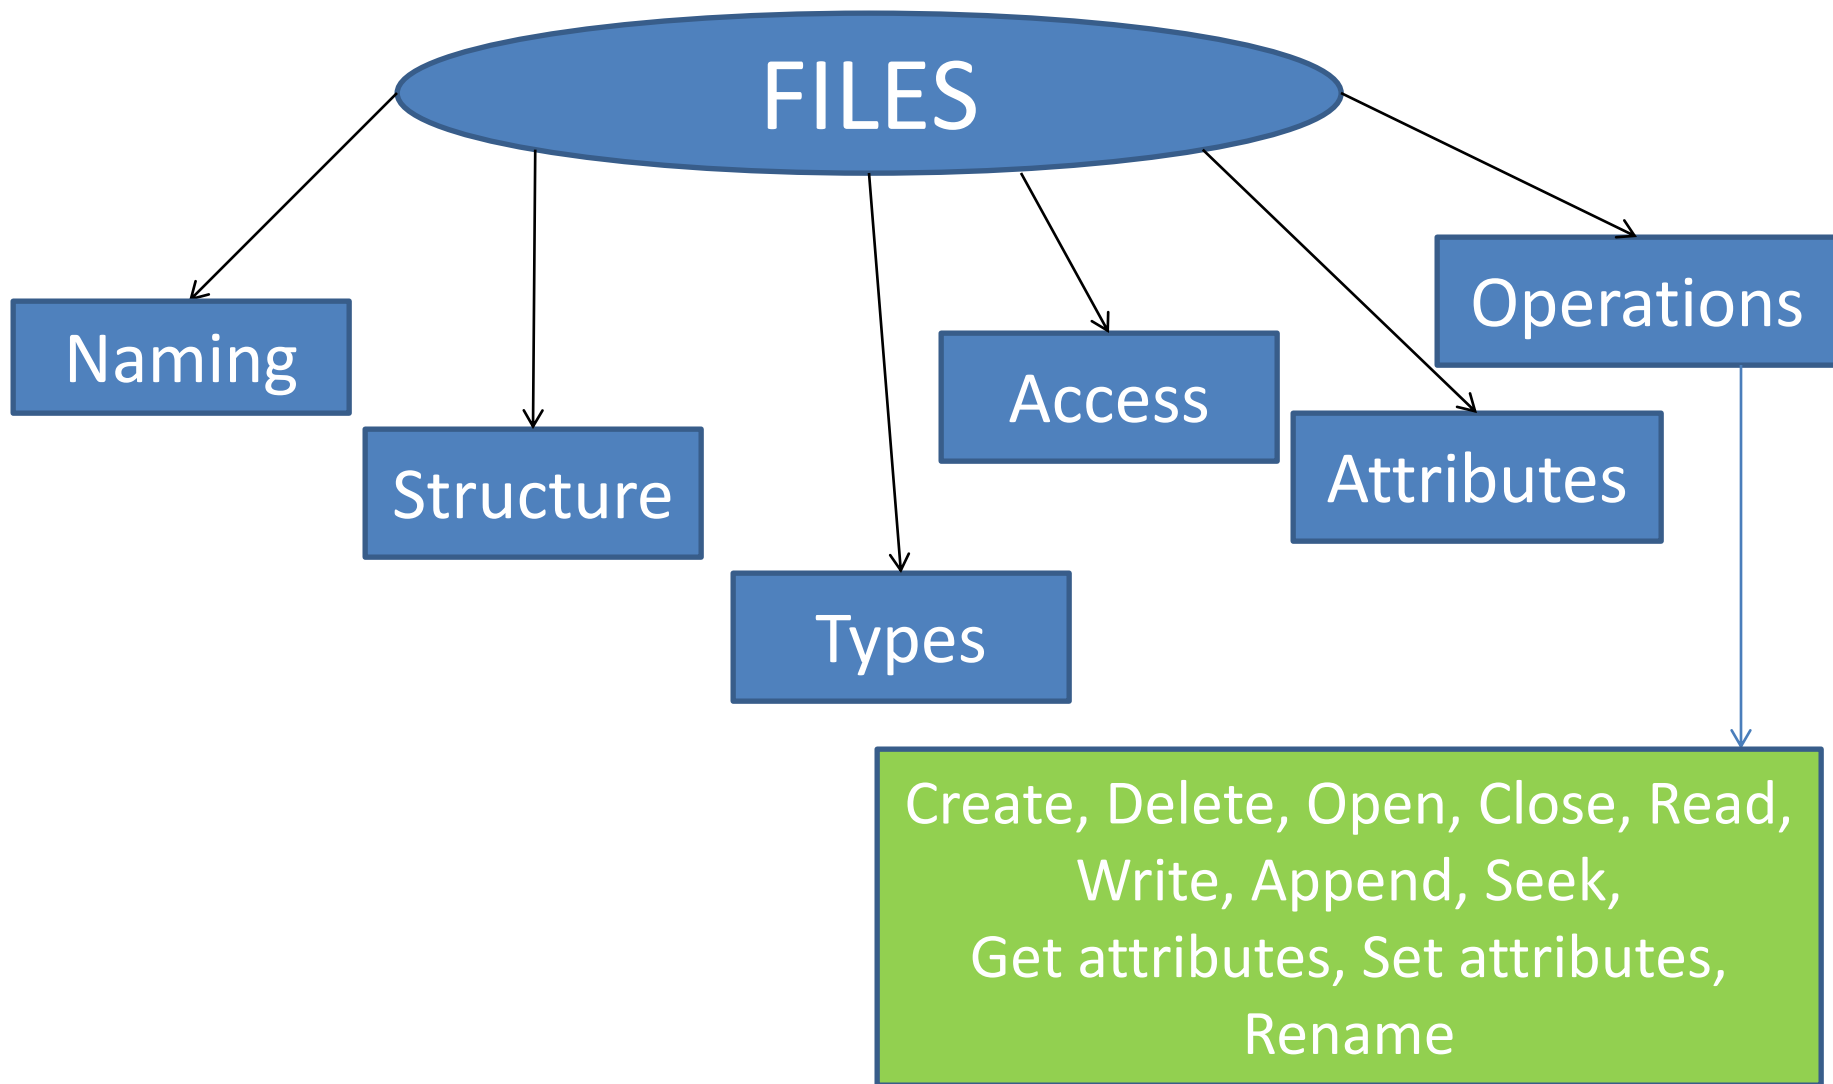


- Regular files
 - ASCII
 - Binary
- Character special
 - to model serial devices (printers, networks, ...)
- Block special
 - to model disks





Attribute	Meaning
Protection	Who can access the file and in what way
Password	Password needed to access the file
Creator	ID of the person who created the file
Owner	Current owner
Read-only flag	0 for read/write; 1 for read only
Hidden flag	0 for normal; 1 for do not display in listings
System flag	0 for normal files; 1 for system file
Archive flag	0 for has been backed up; 1 for needs to be backed up
ASCII/binary flag	0 for ASCII file; 1 for binary file
Random access flag	0 for sequential access only; 1 for random access
Temporary flag	0 for normal; 1 for delete file on process exit
Lock flags	0 for unlocked; nonzero for locked
Record length	Number of bytes in a record
Key position	Offset of the key within each record
Key length	Number of bytes in the key field
Creation time	Date and time the file was created
Time of last access	Date and time the file was last accessed
Time of last change	Date and time the file was last changed
Current size	Number of bytes in the file
Maximum size	Number of bytes the file may grow to



Dealing with Files in C (As an Example)

First step

- Declaration:

```
FILE *fptr1, *fptr2 ;
```

Opening Files

- The statement:

```
fptr1 = fopen ( "filename", "r" );
```

would open the file filename for input (reading).

- r: read
- w: write
- a: append
- ... there are some more

Testing for Successful Open

- If the file was not able to be opened, then the value returned by the *fopen* routine is NULL.
- For example, let's assume that the file *mydata* does not exist. Then:

```
FILE *fptr1 ;  
fptr1 = fopen ( "myfile", "r" );  
if (fptr1 == NULL)  
{  
    printf ("File 'mydata' did not open.\n");  
}
```

Reading From Files

- In the following segment of C language code:

```
int a, b ;  
FILE *fptr1, *fptr2 ;  
fptr1 = fopen ( "mydata", "r" ) ;  
fscanf ( fptr1, "%d%d", &a, &b) ;
```

the *fscanf* function would read values from the file "pointed" to by *fptr1* and assign those values to *a* and *b*.

End of File

- The end-of-file indicator informs the program when there are no more data (no more bytes) to be processed.
- There are a number of ways to test for the end-of-file condition. One is to use the *feof* function which returns a *true* or *false* condition:

```
fscanf (fptr1, "%d", &var) ;  
if ( feof (fptr1) )  
{  
    printf ("End-of-file encountered.\n");  
}
```

- Another (better) way of testing EOF:

```
while(fscanf(fp,"%d ", &current) == 1)  
{  
  
}
```

Writing To Files

```
int a = 5, b = 30;  
FILE *fptr2 ;  
fptr2 = fopen ( "filename", "w" );  
fprintf ( fptr2, "%d %d\n", a, b );
```

the **fprintf** functions would write the values stored in **a** and **b** to the file "pointed" to by **fptr2**.


Closing Files

```
fclose ( fptr1 );
```

Once the files are open, they stay open until you close them or end the program (which will close all files.)

Implementing Files

(What happens at the low-level?)

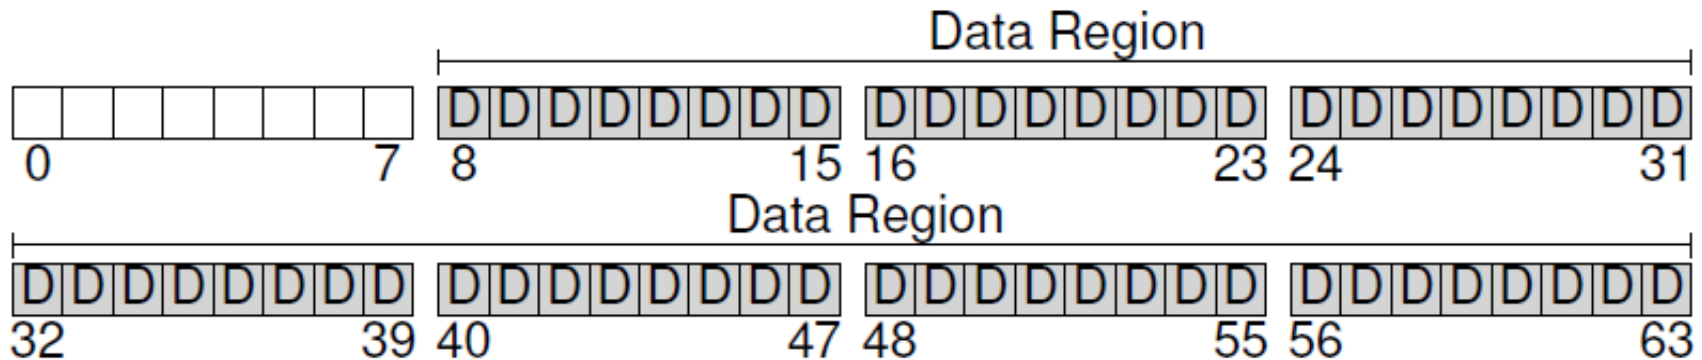


Assumptions:

- A disk is a sequence of blocks.
- All blocks are of the same size.
- If the disk has N blocks, they are addressed $0 \rightarrow N-1$

Implementing Files

(What happens at the low-level?)

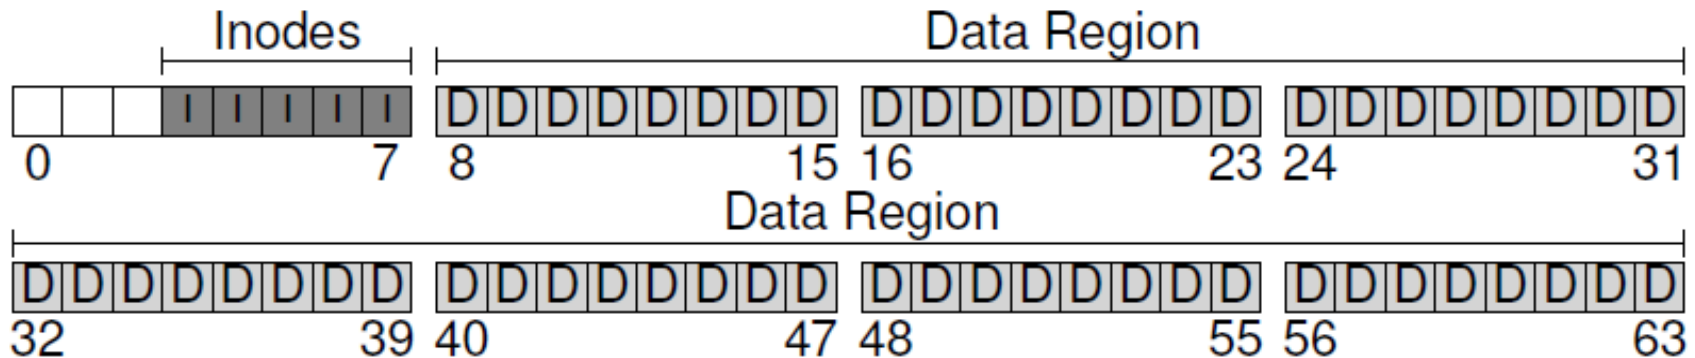


- Most of the disk is reserved for the user data, shown as 'D' in the figure above.
- Each file consists of one or more blocks, not necessarily contiguous.

How to keep track which block(s) belong to which file?

Implementing Files

(What happens at the low-level?)



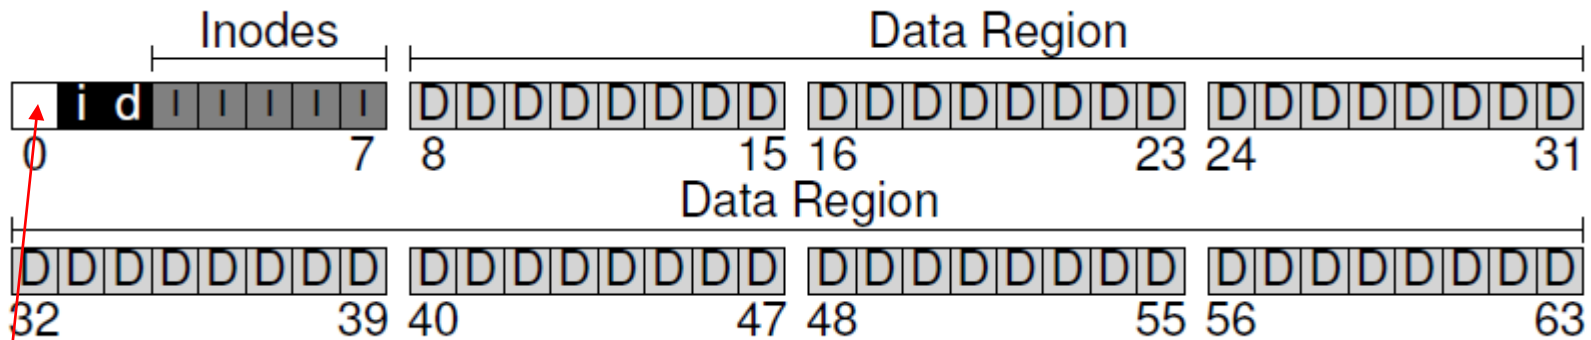
- Each file has an I-node, indicated by the “I” in the figure above.
- I-node stands for “Index-Nodes”, partially for historical reasons.
- I-node contains all the meta-data of the file (security, pointer to first block, etc).
- Each block can hold many i-nodes because an i-node is small (e.g. 256 bytes) and a block is 4KB or more).

How to keep track:

- Which block holds i-nodes and which block holds user data?
- How to keep track of free blocks?

Implementing Files

(What happens at the low-level?)

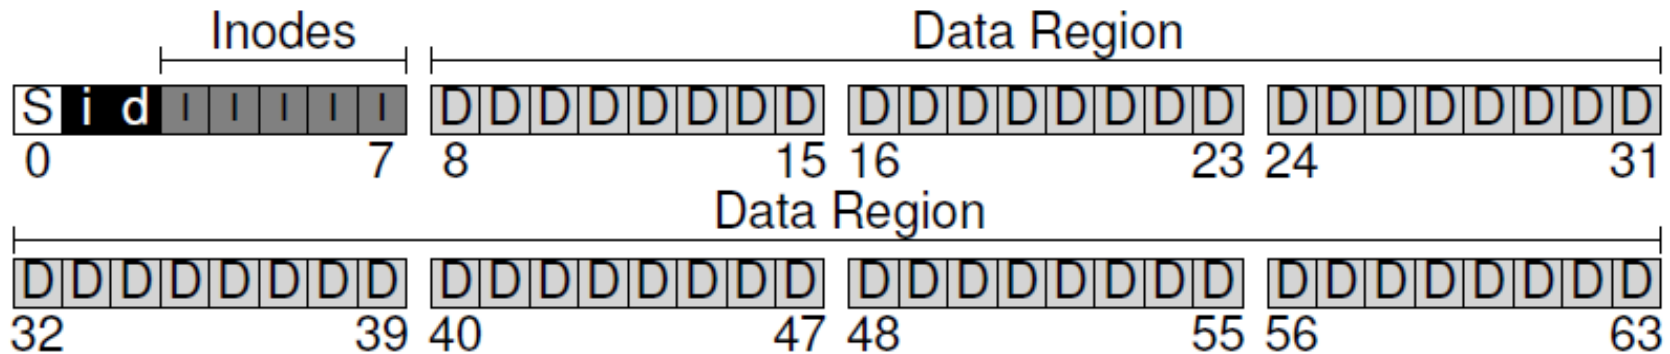


- Both 'i' and 'd' blocks are bitmaps
- A bit can be 0 or 1
- 0 means block is empty and 1 means it is used.
- This way, we keep track of empty and used blocks in both in Inodes region and the data region.

What is the job of that empty block at the beginning?

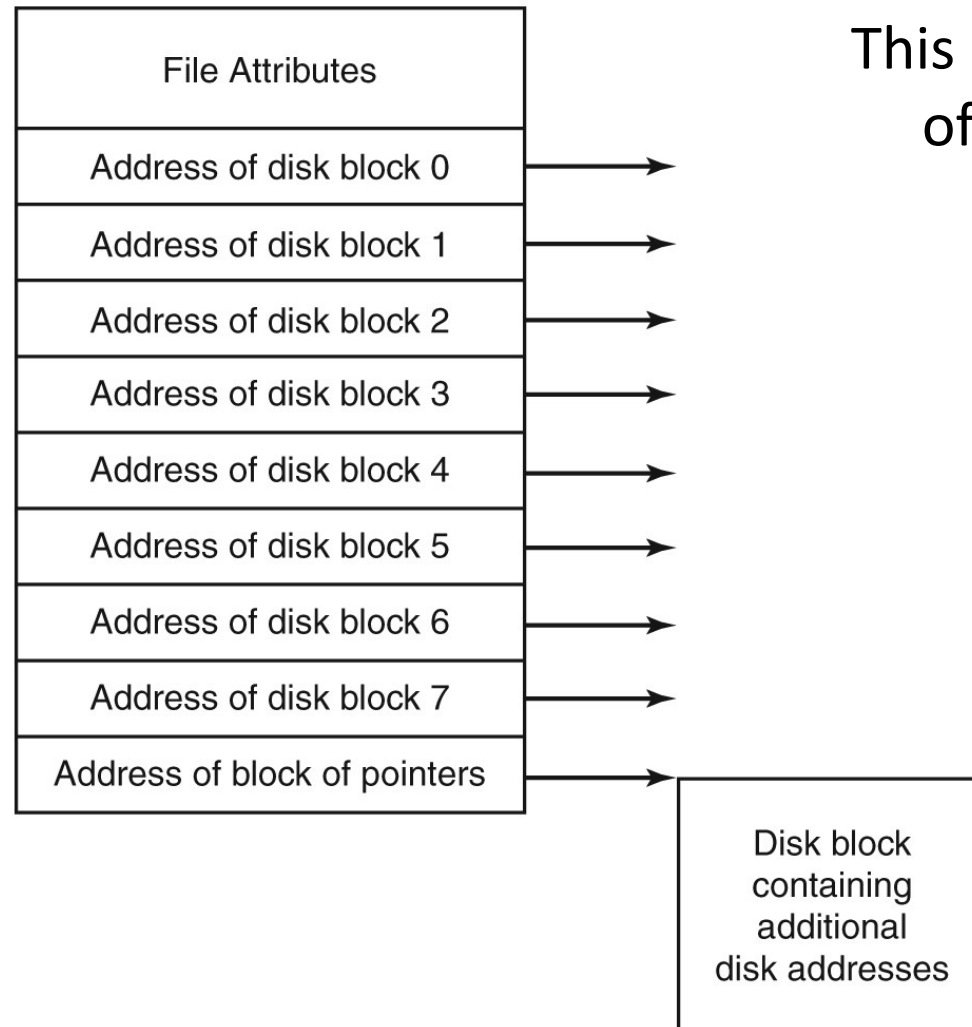
Implementing Files

(What happens at the low-level?)



- S stands for “superblock”.
- Contains information about this file-systems (e.g. type, number of i-nodes, etc).
- The OS reads the S block and attach this volume the file-system tree if there are other volumes.
- A volume can be a disk partition, another disk, etc.
- A directory is just a special type of files.

How to deal with large files?



This is an example of an i-node.

Summary of the Big Picture

- The programmer accesses the file using the API provided by the programming language used.
- This is compiled, and later executed, to a system call.
- Using the device independent I/O layer in the OS, the file access is translated to block(s) access request(s).
 - This is done using the file-system organization we saw in the several previous slides (i.e. I-nodes, etc).
- The device driver takes these requests and translate them to low-level access of the disk (e.g. read sector number a from track b on plate c).

Conclusions

- Files and file system are major parts of an OS.
- We have just scratched the surface of file systems. There are more sophisticated way of implementing it but what we saw is the basics needed.
- Files and File system are the OS way of abstracting storage.