



Operating Systems

Processes - I

Mohamed Zahran (aka Z)


mzahran@cs.nyu.edu

<http://www.mzahran.com>



OS Management of Application Execution

- Resources are made available to multiple applications.
- The CPU(s) is/are switched among multiple **processes** so all will appear to be progressing.
 - This is called **time-sharing**.
- **Main Goal:** The processor, memory, and I/O devices can be used **efficiently**
- In the *processes lectures*, we will concentrate on how to use CPUs efficiently among processes.
- This requires some **mechanisms** and **policies**.



- A set of algorithms
- Make high level decisions:
 - Which process to run next?
 - When to stop a process?
 - ...
- Decision is based on some criteria:
 - Performance
 - Power
 - Priority
 -
- Low-level protocol to implement a functionality
- Example:
 - How to switch one process for the other (context-switching)
 - Collecting process's metrics



```
graph LR; A([Executable File]) --> B([CPU executes the code]);
```

Executable File

CPU
executes the code

Sometime called:

- Object file
- Binary file

Initially, there is an executable file of a program residing on your disk.

A command is issued to the OS to start executing this program.

The loader, part of the OS, extracts needed info from the executable file, loads them in memory, the OS does some initialization, and the CPU is told to start executing.

When the CPU begins to execute the program code, we refer to this executing entity as a ***process***

Three Kinds of Object Files (sometimes called modules)

- **Relocatable object file** (.o file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source (.c) file
- **Executable object file** (a.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- **Shared object file** (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- Standard binary format for object files
 - Originally proposed by AT&T System V Unix, later adopted by BSD Unix variants and Linux
- One unified format for
 - Relocatable object files (`.o`),
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- Generic name: ELF binaries

ELF Object File Format

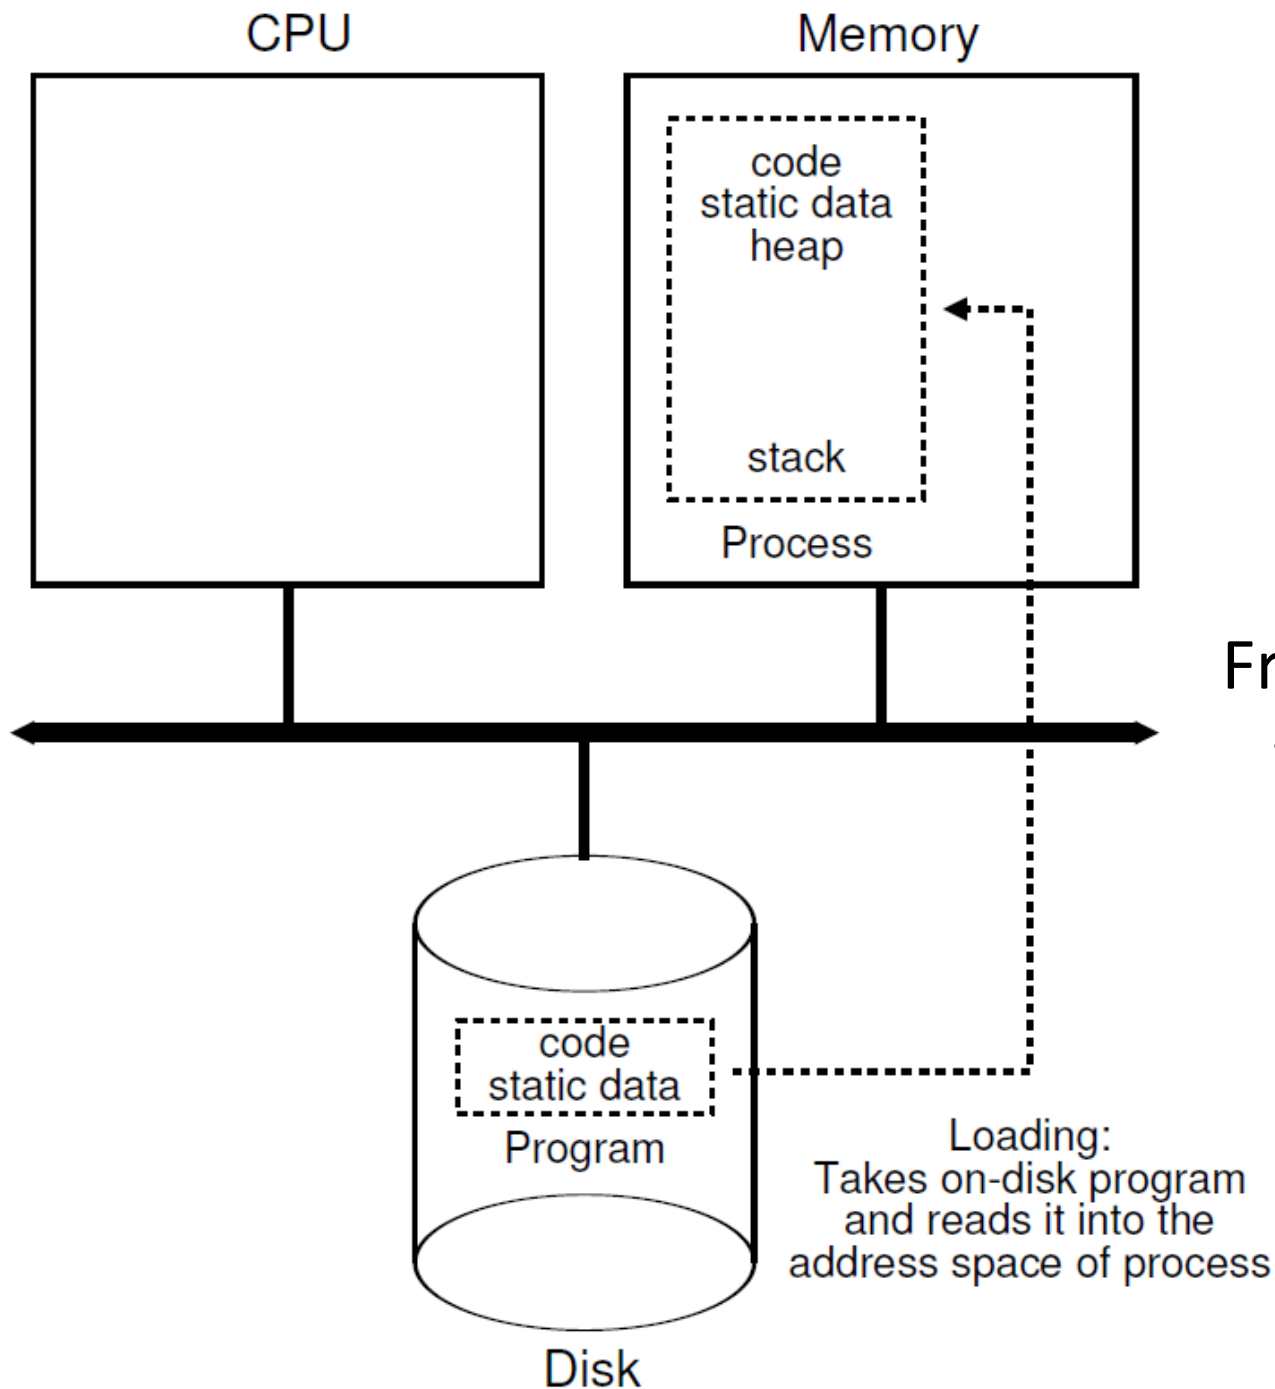
- Elf header
 - Word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- Segment header table
 - Page size, virtual addresses memory segments (sections), segment sizes.
- .text section
 - Code
- .rodata section
 - Read only data: jump tables, ...
- .data section
 - Initialized global variables
- .bss section (**B**lock **S**tarted by **S**ymbol)
 - Uninitialized global variables
 - Only the length but no data
 - Later, the program loader will allocate memory for it.

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

ELF Object File Format (cont.)

- **.symtab section**
 - Symbol table
 - Procedure and global variable names
- **.rel.text section**
 - Relocation info for **.text** section
 - Addresses of instructions that will need to be modified in the executable
- **.rel.data section**
 - Relocation info for **.data** section
 - Addresses of pointer data that will need to be modified in the merged executable
- **.debug section**
 - Info for symbolic debugging (`gcc -g`)
- **Section header table**
 - Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table



From a program
to a process.

The Process Model

A process is an instance of an executing program and includes

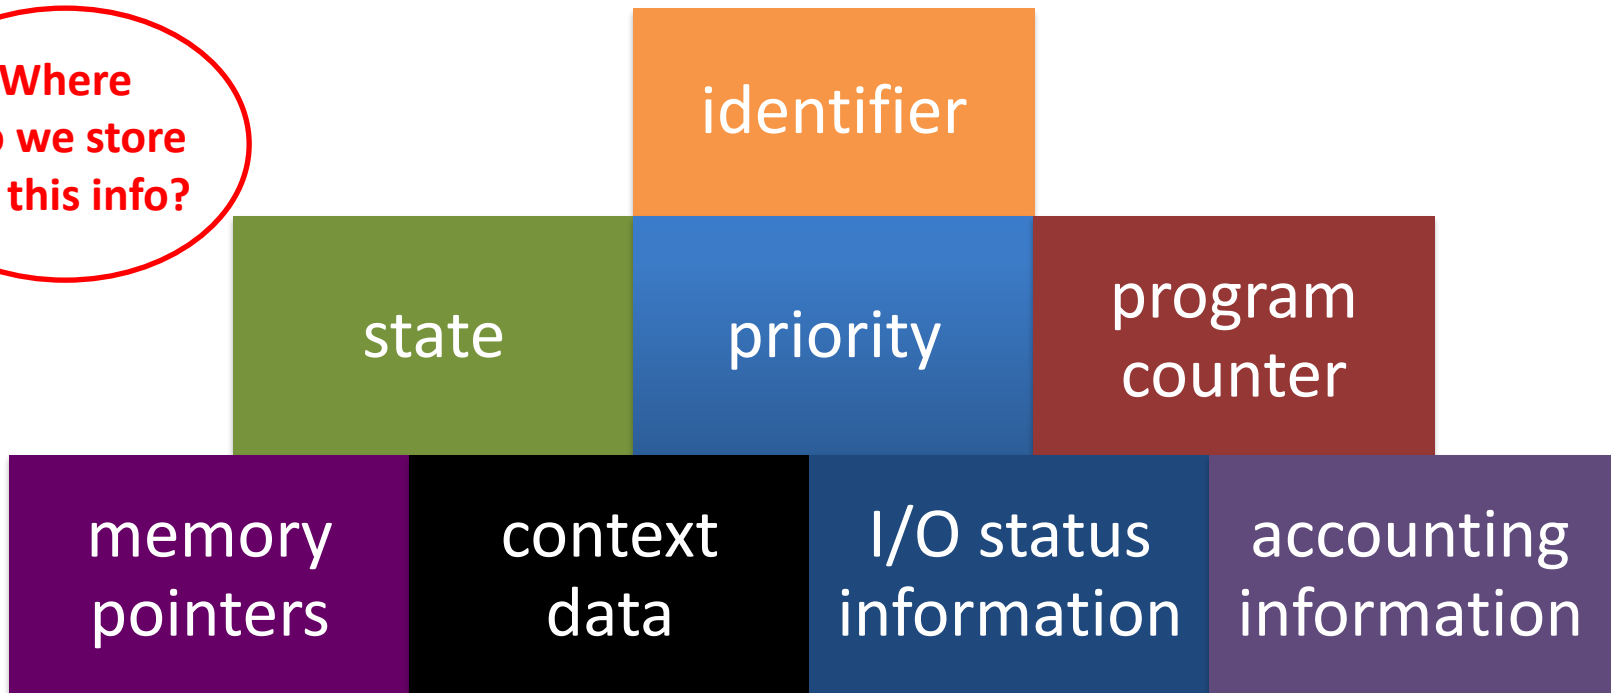
- Program counter
- Registers
- Variables
- Code
- ...

If a program is running twice, does it count as two processes? or one?

Process

- While the program is executing, its process can be uniquely characterized by a number of elements, including:

Where
do we store
all this info?

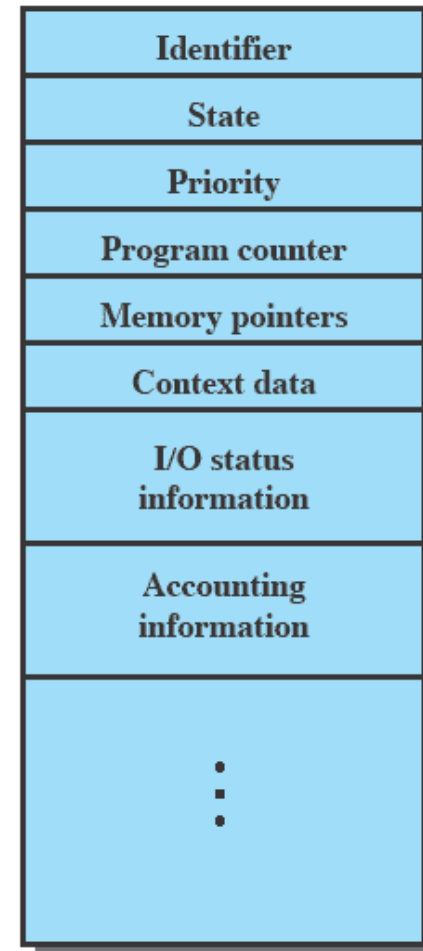


Data Structure

- The OS keeps a list of all processes in the system. → **process list** (also called process table or control table).
- Each entry of this list contains information about a process. → **process control block**

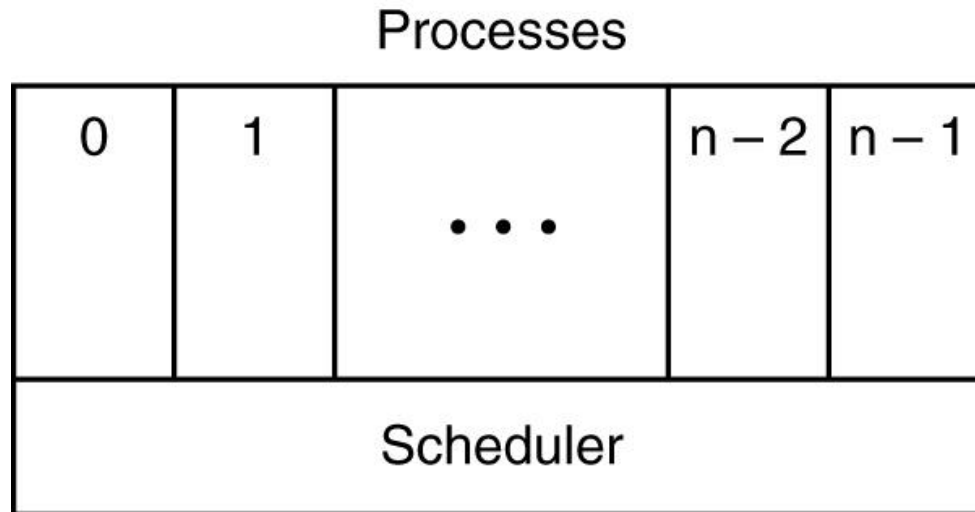
Process Control Block

- Contains the process elements
- It is possible to **interrupt** a running process and later resume execution as if the interruption had not occurred
- Created and managed by the operating system
- Key tool that allows support for multiple processes




Multiprogramming

- One CPU and several processes
- CPU switches from process to process quickly




One program counter




(a)

Four program counters



(b)




What Really Happens

What We Think It Happens


If we run the same program several times,
will we get the same execution time?

One program counter




(a)

Four program counters




(b)



What Really Happens

What We Think It Happens

Example



time

1 5000
2 5001
3 5002
4 5003
5 5004
6 5005

----- Timeout

7 100
8 101
9 102
10 103
11 104
12 105

13 8000
14 8001
15 8002
16 8003

----- I/O Request

17 100
18 101
19 102
20 103
21 104
22 105

23 12000
24 12001
25 12002
26 12003

Instruction address

27 12004
28 12005

----- Timeout

29 100
30 101
31 102
32 103
33 104
34 105

35 5006
36 5007
37 5008
38 5009
39 5010
40 5011

----- Timeout

41 100
42 101
43 102
44 103
45 104
46 105

47 12006
48 12007
49 12008
50 12009
51 12010
52 12011

----- Timeout

Small program that switches the processor from one process to another (also called Scheduler)

Operations with Processes (i.e. APIs)

- Creation
- Termination or destroying
- Status
- Wait
- Misc: suspend, resume, inquire, ...

Process Creation:

When Does it Happen?

- System initialization
 - At boot time
 - Foreground
 - Background (daemons)
- Execution of a process creation system call by a running process
- A user request
- A batch job
- Created by OS to provide a service
- Interactive logon

Process Termination: When Does it Happen?

- Normal exit (voluntary)
- Error exit (voluntary)
- Fatal error (involuntary)
- Killed by another process (involuntary)

Process Termination: More Scenarios


Normal completion	The process executes an OS service call to indicate that it has completed running.
Time limit exceeded	The process has run longer than the specified total time limit. There are a number of possibilities for the type of time that is measured. These include total elapsed time ("wall clock time"), amount of time spent executing, and, in the case of an interactive process, the amount of time since the user last provided any input.
Memory unavailable	The process requires more memory than the system can provide.
Bounds violation	The process tries to access a memory location that it is not allowed to access.
Protection error	The process attempts to use a resource such as a file that it is not allowed to use, or it tries to use it in an improper fashion, such as writing to a read-only file.
Arithmetic error	The process tries a prohibited computation, such as division by zero, or tries to store numbers larger than the hardware can accommodate.
Time overrun	The process has waited longer than a specified maximum for a certain event to occur.
I/O failure	An error occurs during input or output, such as inability to find a file, failure to read or write after a specified maximum number of tries (when, for example, a defective area is encountered on a tape), or invalid operation (such as reading from the line printer).
Invalid instruction	The process attempts to execute a nonexistent instruction (often a result of branching into a data area and attempting to execute the data).
Privileged instruction	The process attempts to use an instruction reserved for the operating system.
Data misuse	A piece of data is of the wrong type or is not initialized.
Operator or OS intervention	For some reason, the operator or the operating system has terminated the process (e.g., if a deadlock exists).
Parent termination	When a parent terminates, the operating system may automatically terminate all of the offspring of that parent.
Parent request	A parent process typically has the authority to terminate any of its offspring.

Process State


- Depending on the implementation, there can be several possible **state models**.
- Each state model shows:
 - number of states
 - type of each state
 - when to move from a state to another

Process State: Two-State Model

The simplest model

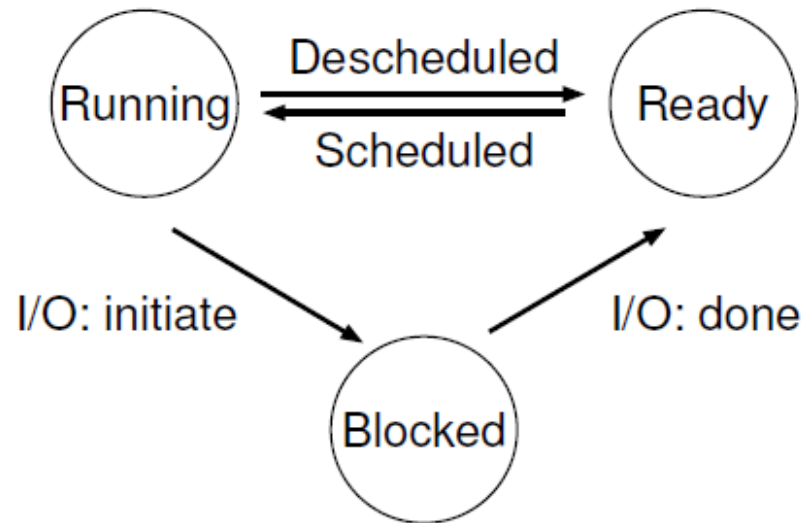


Process State: Three-State Model



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

Process State: Three-State Model



Summarizing the previous slide

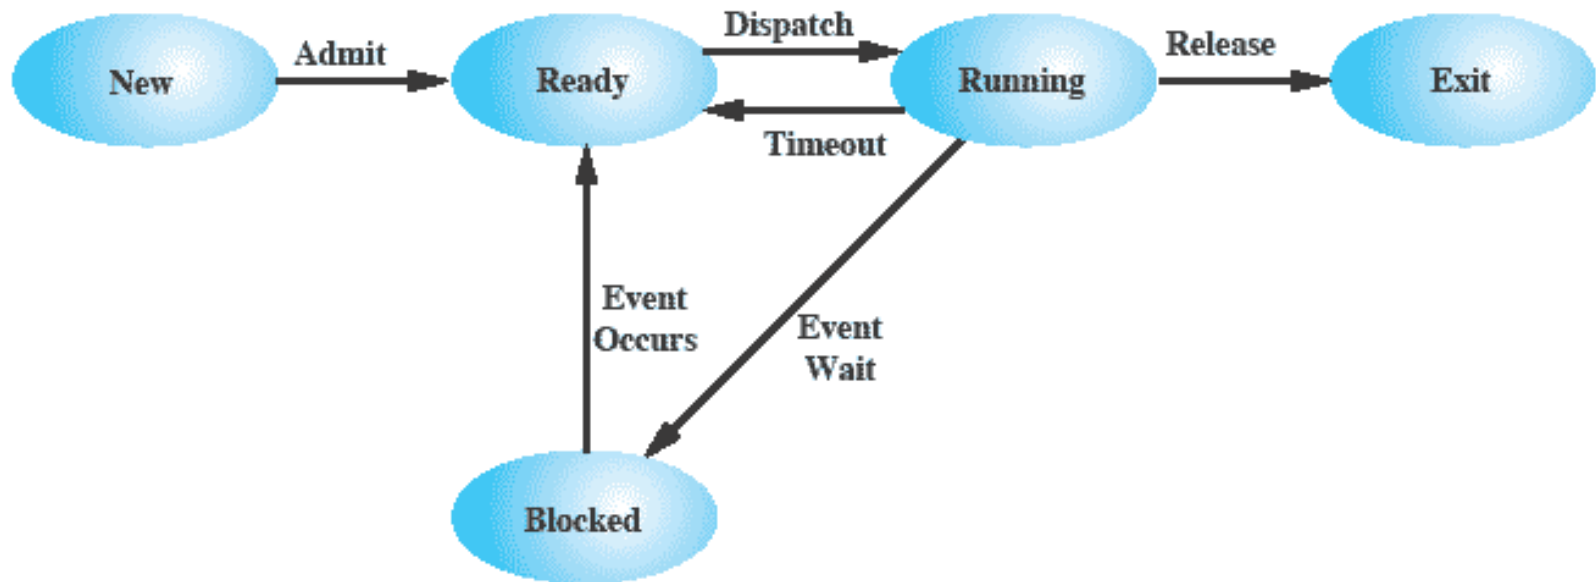
Example:

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process ₀ now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process ₁ now done

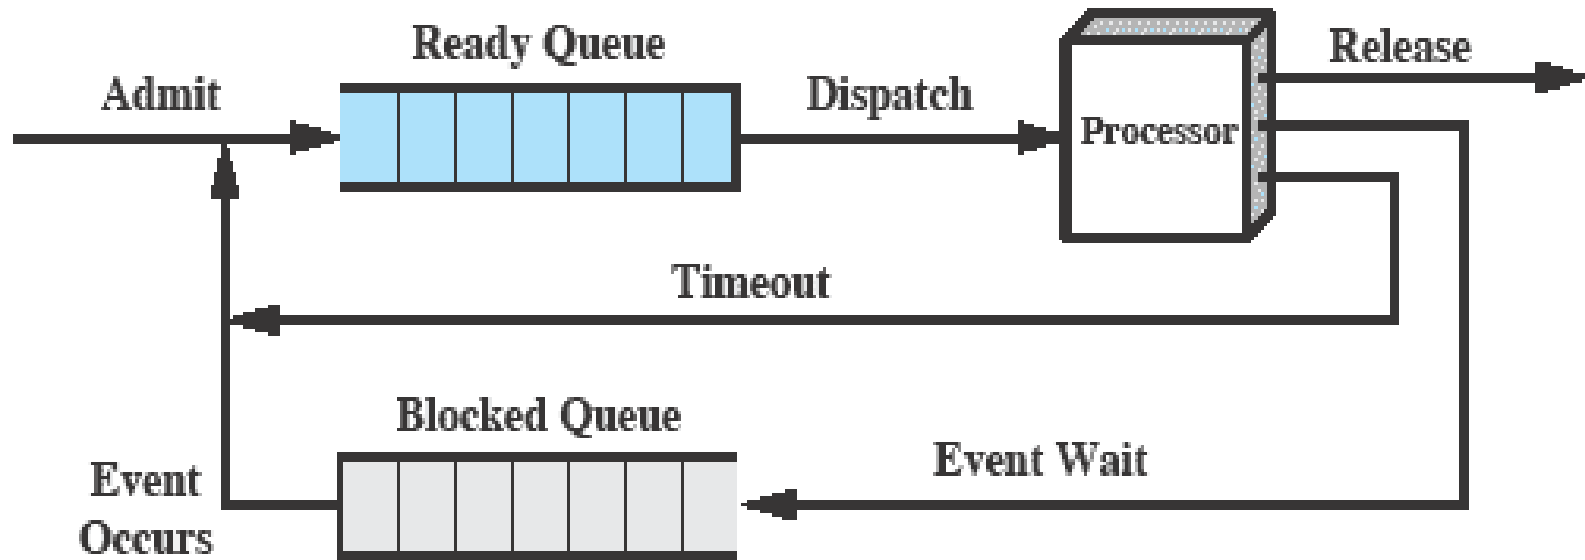
Example:

Time	Process ₀	Process ₁	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process ₀ initiates I/O
4	Blocked	Running	Process ₀ is blocked,
5	Blocked	Running	so Process ₁ runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process ₁ now done
9	Running	–	
10	Running	–	Process ₀ now done

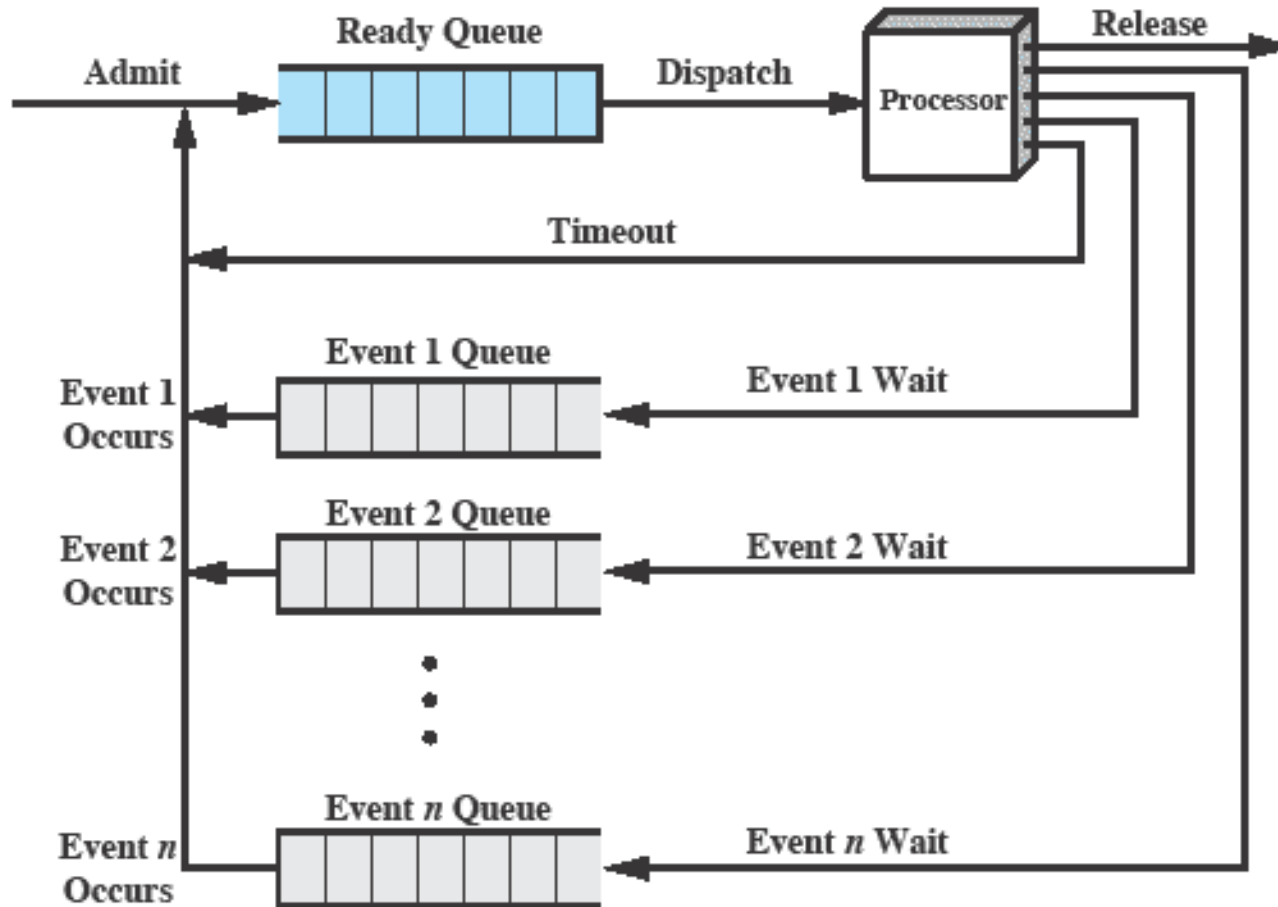
Process State Five-State Model




Using Queues to Manage Processes



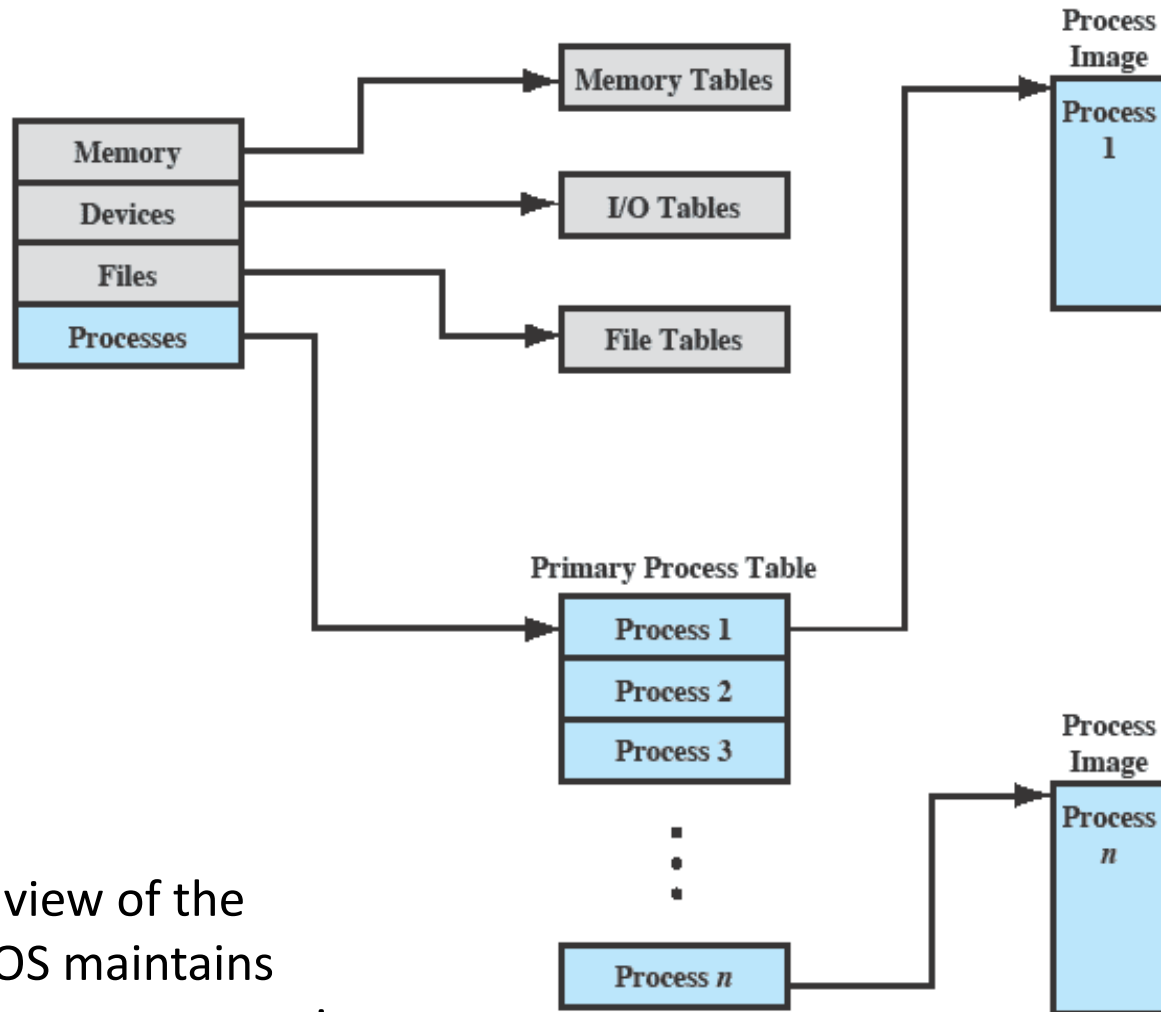
Using Queues to Manage Processes



One Extra State!



Swapped to disk



Conceptual view of the tables that OS maintains in order to manage execution of processes on resources.

A Bit About Interrupts

- **Interrupt means:**
 - Current running process is suspended
 - OS takes control (i.e. the machine moves from user mode to kernel mode)
- **Interrupt occurs due to many scenarios, for example:**
 - Time out of current running process
 - Hardware interrupt from an I/O device
 - Page fault

A Bit About Interrupts


1. Hardware saves program counter, registers, etc of the current process.
2. Hardware loads program counter of the interrupt service routine (ISR).
3. ISR runs.
4. When done, the next process to run is picked.
5. Program counter, registers, etc of the picked process are loaded.
6. The picked process starts running,

Moving from User Mode to Kernel Mode

- Interrupts
- System call
- Exception (e.g. divide by zero, segmentation fault, etc)


Simple Modeling of Multiprogramming

- A process spends fraction p waiting for I/O
- Assume n processes in memory at once
- The probability that all processes are waiting for I/O at once is p^n
- So -> **CPU Utilization = $1 - p^n$**




Multiprogramming lets processes use the CPU when it would otherwise become idle.


Executing the OS Itself



(a) Separate kernel



(b) OS functions execute within user processes



(c) OS functions execute as separate processes

Conclusions

- Process is the most central concept in OS
- Process is the main way by which OS virtualizes the CPU
- Even with many cores (i.e., many CPUs), we still need multiprogramming (i.e., time-sharing).