



Operating Systems

Processes - II

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



Two Main Challenges

- How can OS implement virtualization of the CPU without adding excessive overhead to the system?
- How can the OS run processes efficiently while retaining control over the CPU?

Straightforward Solution: Direct Execution

OS

Create entry for process list
Allocate memory for program
Load program into memory
Set up stack with argc/argv
Clear registers
Execute **call** main()

Free memory of process
Remove from process list

Program

Q1: How does the OS guarantee
that this program won't do
anything nasty?



Run main()
Execute **return** from main




Q2: How can the OS stop this process from
running and switch to another process?

Answer to Q1: Introducing Kernel mode and User mode

OS @ boot
(kernel mode)
initialize trap table

Hardware

remember address of...
syscall handler



OS @ run
(kernel mode)

Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from-trap

Hardware

restore regs
(from kernel stack)
move to user mode
jump to main

Program
(user mode)

Run main()
...
Call system call
trap into OS

save regs
(to kernel stack)
move to kernel mode
jump to trap handler

Handle trap
Do work of syscall
return-from-trap

restore regs
(from kernel stack)
move to user mode
jump to PC after trap

...
return from main
trap (via `exit()`)

Free memory of process
Remove from process list

**Trap table contains the addresses
of the different trap handles
(aka Interrupt Service Routines)**

What Happened?

- Processes run at user mode so cannot access the hardware.
- If a process needs something from the hardware, it makes a **system call** with a specific **system call number**.
- This switches the machine to kernel mode and the OS starts executing the service routine specified by the number.
- Now, a process has limited power over the system and cannot do something that we don't want it to do.

Answer to Q2: We need help from the hardware

- If a process is using the CPU, this means the OS is not running (you can scale that to multi-processes with multicore).
- If so, how can an OS stop a process to switch to another process?
- Solution: **timer interrupts**

Timer Interrupt

- A device programmed to raise and interrupt every x milliseconds
 - x is set by the OS
 - OS starts the timer at boot time.
- When the interrupt occurs, the machine switches to kernel and the OS takes control.
 - The **scheduler** part of the OS
 - The scheduler decides: continue running the current process or **context switch** to another one.

APIs for dealing with Processes in C

Don't forget to include the
following header files:

```
#include<unistd.h>
```

```
#include<sys/types.h>
```

```
#include<sys/wait.h>
```

Basic Syscalls for Managing Processes

- **fork** spawns new process
 - Called once, returns twice
- **exit** terminates own process
 - Puts it into "zombie" status until its parent reaps
- **wait** and **waitpid** wait for and reap terminated children
- **execve** runs new program in existing process
 - Called once, never returns

fork: Creating New Processes

- `int fork(void)`
 - creates a new process (child process) that is identical to the calling process (parent process)
- Fork is called *once* but returns *twice*

```
pid_t pid = fork();  
if (pid == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent: child pid is %d\n", pid);  
}
```

Return 0 to the child process

Return child's pid to the parent

Understanding fork

Process n



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

Child Process m



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



pid = m

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



pid = 0

```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```



```
pid_t pid = fork();
if (pid == 0) {
    printf("hello from child\n");
} else {
    printf("hello from parent\n");
}
```

hello from parent

hello from child

Fork Example #1


- Parent and child both run same code
 - Distinguish parent from child by return value from `fork`
- Start with same state, but each has private copy of memory
 - Including shared output file descriptor
 - Relative ordering of their print statements undefined

```
void fork1()
{
    int x = 1;
    pid_t pid = fork();
    if (pid == 0) {
        printf("Child has x = %d\n", ++x);
    } else {
        printf("Parent has x = %d\n", --x);
    }
    printf("Bye from process %d with x = %d\n", getpid(), x);
}
```

Fork Example #2

- Both parent and child can continue forking


```
void fork2()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("Bye\n");  
}
```



Fork Example #3

- Both parent and child can continue forking


```
void fork3()  
{  
    printf("L0\n");  
    fork();  
    printf("L1\n");  
    fork();  
    printf("L2\n");  
    fork();  
    printf("Bye\n");  
}
```



Fork Example #4

- Both parent and child can continue forking


```
void fork4()  
{  
    printf("L0\n");  
    if (fork() != 0) {  
        printf("L1\n");  
        if (fork() != 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



Fork Example #5

- Both parent and child can continue forking

```
void fork5()  
{  
    printf("L0\n");  
    if (fork() == 0) {  
        printf("L1\n");  
        if (fork() == 0) {  
            printf("L2\n");  
            fork();  
        }  
    }  
    printf("Bye\n");  
}
```



exit: Ending a process

- `void exit(int status)`
 - exits a process
 - Normally return with status 0
 - `atexit(function_name)` make `function_name` execute upon exit

```
void cleanup(void) {  
    printf("cleaning up\n");  
}  
  
void fork6() {  
    atexit(cleanup);  
    fork();  
    exit(0);  
}
```

Zombies!

- Idea
 - When process terminates , still consumes system resources (i.e. an entry in process table)
 - Why? So that parents can learn of children's exit status
 - Called a "zombie"
- Reaping
 - Performed by parent on terminated child
 - Parent is given exit status information
 - OS discards process
- What if parent doesn't reap?
 - If parent has terminated, then child will be reaped by **init process** (the great-great-...-grandparent of all user-level processes)

Zombie Example

```
linux> ./forks7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6639 ttyp9        00:00:03 forks
 6640 ttyp9        00:00:00 forks <defunct>
 6641 ttyp9        00:00:00 ps
linux> kill 6639
[1]      Terminated
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6642 ttyp9        00:00:00 ps
```


```
void fork7()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
            getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
            getpid());
        while (1) ; /* Infinite loop */
    }
}
```

- `ps` shows child process as "defunct"
- Killing parent allows child to be reaped by `init`

wait: Synchronizing with Children

- `int wait(int *child_status)`
 - Blocks until some child exits, return value is the `pid` of terminated child
 - If multiple children completed, will take in arbitrary order (use `waitpid` to wait for a specific child)

```
void fork8() {  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
    }  
    else {  
        printf("HP: hello from parent\n");  
        wait(NULL);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
    exit(0);  
}
```



This is how child process is reaped by parent process.

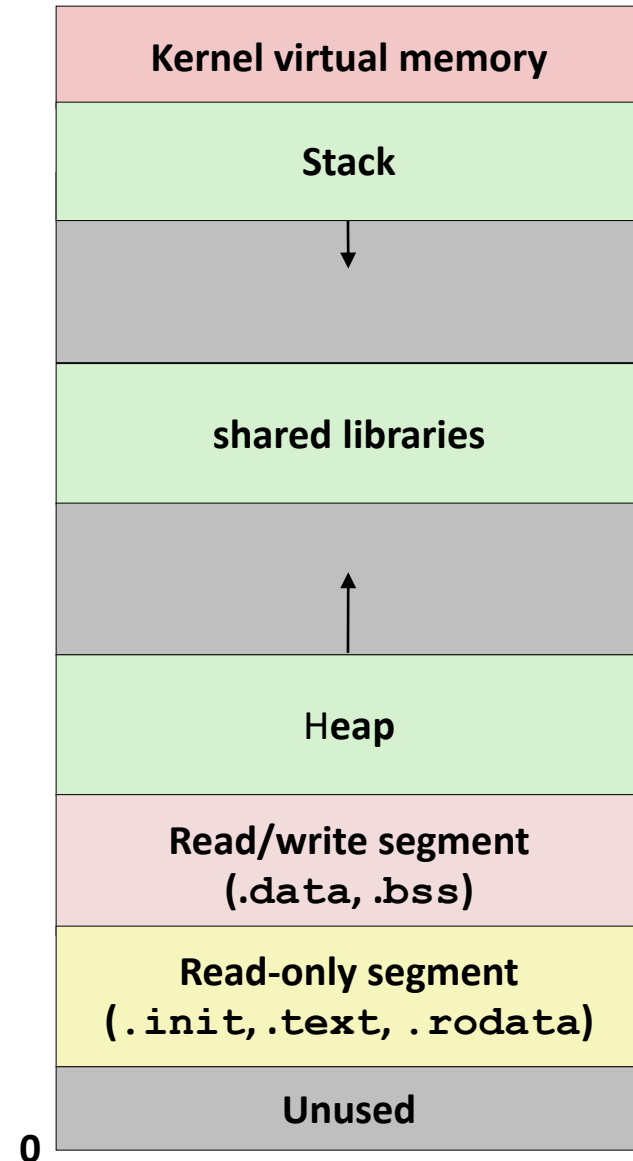
execve

- **int execve(char *fname, char *argv[], char *envp[])**
 - Executes program named by `fname`

```
if ((pid = fork()) == 0) { /* Child runs user job */
    if (execve(argv[0], argv, environ) < 0) {
        printf("%s: Command not found.\n", argv[0]);
        exit(0);
    }
}
```

`execve`: Load a new program image

- `execve` causes OS to overwrite code, data, and stack of process
 - keeps pid and open files



Conclusions

- OS must ensure that a running process will not cause any harm to the whole system.
- OS runs in kernel mode and has access to protected instructions (i.e., a set of machine code instructions that cannot be executed while the machine is in user mode).